

## ASSESSMENT 2

ASSESSMENT DATE	26-09-2022
STUDENT NAME	Santhavaliyan.m
STUDENT ROLL NUMBER	713119205007
MAXIMUM MARKS	2 Marks

### 1. Download the dataset

### 2. Load the dataset

#### Solution:

```
import pandas as pd
import numpy as np
df=pd.read_csv("C:\\Users\\PC\\Desktop\\Churn_Modelling.csv")
df
```

#### output:

Out[7]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93828.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	93828.63
9996	9997	15596992	Johnstone	516	France	Male	35	10	57369.91	1	1	1	101348.88
9997	9998	15594532	Liu	709	France	Female	35	7	0.00	1	0	1	43761.12
9998	9999	15692355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	93828.63
9999	10000	15596140	Mellors	700	France	Female	39	4	125445.70	1	1	0	113931.57

df.head()

Out[8]:

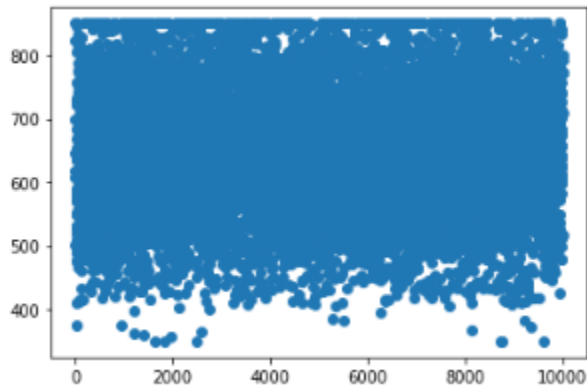
	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93828.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

### 3.perform following operations

#### ➤ univariate analysis

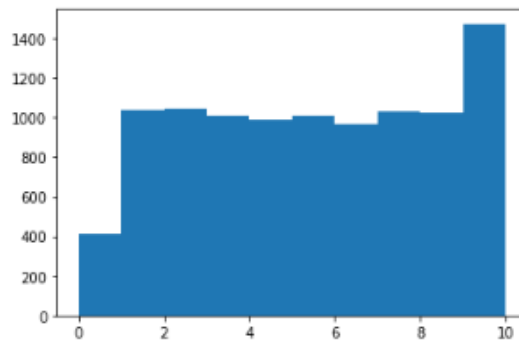
```
import matplotlib.pyplot as plt
import seaborn as sns
plt.scatter(df.index,df['CreditScore'])
plt.show()
```

output:




```
plt.hist(df['Tenure'])
```

```
Out[12]: (array([ 413., 1035., 1048., 1009.,  989., 1012.,  967., 1028., 1025.,
        1474.]),
         array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.]),
         <a list of 10 Patch objects>)
```



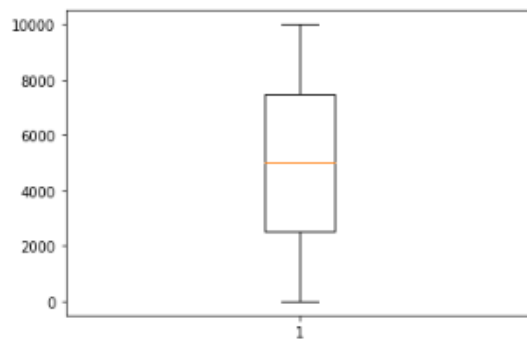
```
sns.distplot(df['NumOfProducts'])
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x25dff3899c8>
```

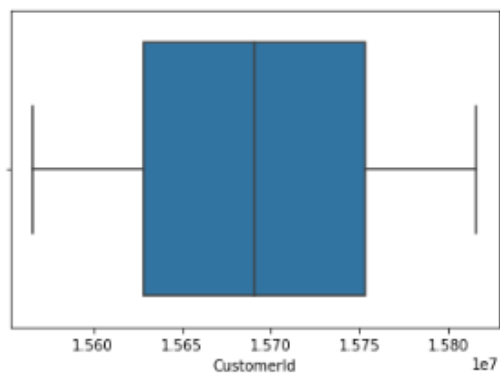


```
plt.boxplot(df['RowNumber'])
```

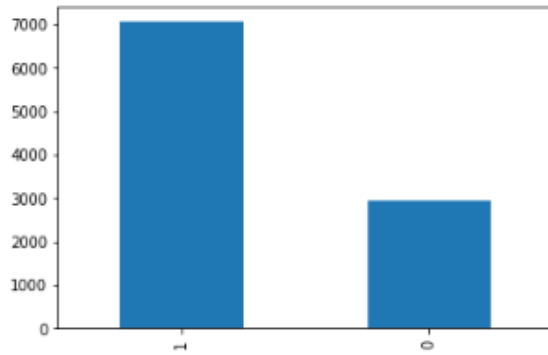
```
plt.show()
```



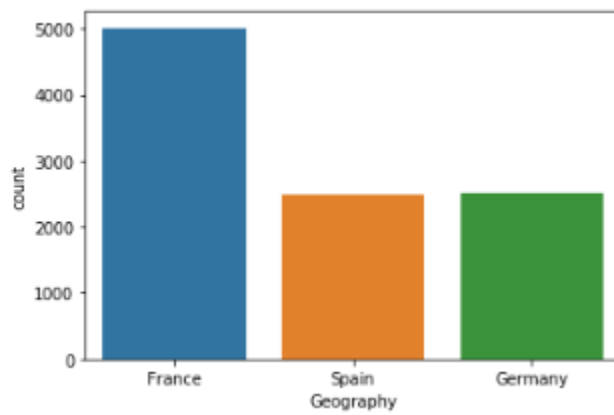
```
sns.boxplot(df['CustomerId'])
```



```
df['HasCrCard'].value_counts().plot.bar()
```



```
sns.countplot(df['Geography'])
```

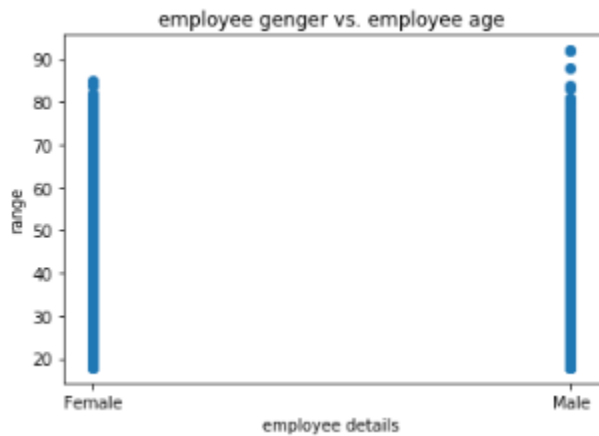


➤ **Bivariate analysis**

```
plt.scatter(df.Gender, df.Age)
```

```
plt.title('employee genger vs. employee age')
```

```
plt.xlabel('employee details') plt.ylabel('range')
```



```
plt.scatter(df.NumOfProducts, df.Balance)
```

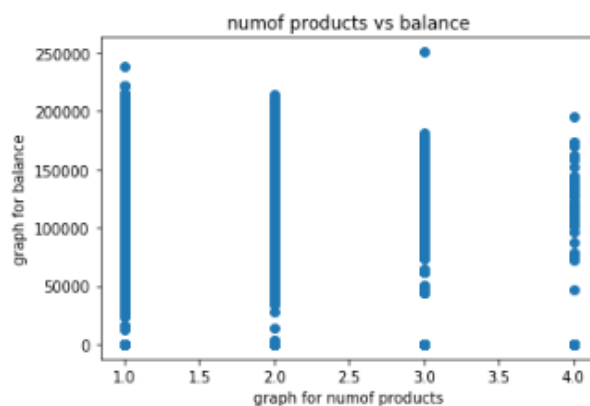
```
plt.title('numof products vs balance')
```

```
plt.xlabel('graph for numof products')
```

```
plt.ylabel('graph for balance')
```

---

```
Out[20]: Text(0, 0.5, 'graph for balance')
```



## ➤ Multivariate analysis

seaborn.pairplot(df)

plt.show()



## 4.descriptive function

df.describe()

Out[6]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500
max	10000.00000	1.581569e+07	850.000000	62.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000

## 5.handle the missing data

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column              Non-Null Count  Dtype
---  -
0   RowNumber            10000 non-null  int64
1   CustomerId           10000 non-null  int64
2   Surname              10000 non-null  object
3   CreditScore          10000 non-null  int64
4   Geography            10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard            10000 non-null  int64
11  IsActiveMember       10000 non-null  int64
12  EstimatedSalary      10000 non-null  float64
13  Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

df.isnull()

Out[14]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSal
0	False	False	False	False	False	False	False	False	False	False	False	False	Fa
1	False	False	False	False	False	False	False	False	False	False	False	False	Fa
2	False	False	False	False	False	False	False	False	False	False	False	False	Fa
3	False	False	False	False	False	False	False	False	False	False	False	False	Fa
4	False	False	False	False	False	False	False	False	False	False	False	False	Fa
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9996	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9997	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9998	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9999	False	False	False	False	False	False	False	False	False	False	False	False	Fa

10000 rows x 14 columns

df.notnull()

Out[15]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	True	True	True	True	True	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	True	True	True	True	True	True
2	True	True	True	True	True	True	True	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True	True	True	True	True	True
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	True	True	True	True	True	True	True	True	True	True	True	True	True
9996	True	True	True	True	True	True	True	True	True	True	True	True	True
9997	True	True	True	True	True	True	True	True	True	True	True	True	True
9998	True	True	True	True	True	True	True	True	True	True	True	True	True
9999	True	True	True	True	True	True	True	True	True	True	True	True	True

10000 rows x 14 columns

df.fillna(0)

Out[16]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

10000 rows x 14 columns

df["Gender"].fillna("No Gender", inplace = True)

df

Out[21]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

10000 rows x 14 columns



```
df.drop("RowNumber",axis=1,inplace=True)
```

```
df
```

```
Out[28]:
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	15619304	Onio	602	France	Female	42	8	159860.80	3	1	0	113931.57
3	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10
...	...	...	...	...	...	...	...	...	...	...	...	...
9995	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	98270.64
9996	15569892	Johnstone	616	France	Male	35	10	57369.61	1	1	1	101699.77
9997	15584532	Liu	709	France	Female	38	7	0.00	1	0	1	42085.58
9998	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	92888.52
9999	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.78

10000 rows x 12 columns

```
print(df.isnull().sum())
```

```
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
dtype: int64
```

```
updated_df = df.dropna(axis=1)
```

```
updated_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   CustomerId      10000 non-null  int64
1   Surname         10000 non-null  object
2   CreditScore     10000 non-null  int64
3   Geography       10000 non-null  object
4   Gender          10000 non-null  object
5   Age             10000 non-null  int64
6   Tenure          10000 non-null  int64
7   Balance         10000 non-null  float64
8   NumOfProducts  10000 non-null  int64
9   HasCrCard       10000 non-null  int64
10  IsActiveMember  10000 non-null  int64
11  EstimatedSalary 10000 non-null  float64
dtypes: float64(2), int64(7), object(3)
memory usage: 937.6+ KB
```

## 6.Finding outliers and replace

```
Q1 = df.quantile(0.25)
```

```
Q3 = df.quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
print(IQR)
```

```
RowNumber      4999.5000
CustomerId      124705.5000
CreditScore      134.0000
Age              12.0000
Tenure           4.0000
Balance          127644.2400
NumOfProducts    1.0000
HasCrCard         1.0000
IsActiveMember    1.0000
EstimatedSalary  98386.1375
Exited           0.0000
dtype: float64
```

```
print(df < (Q1 - 1.5 * IQR))
```

```
(df > (Q3 + 1.5 * IQR))
```

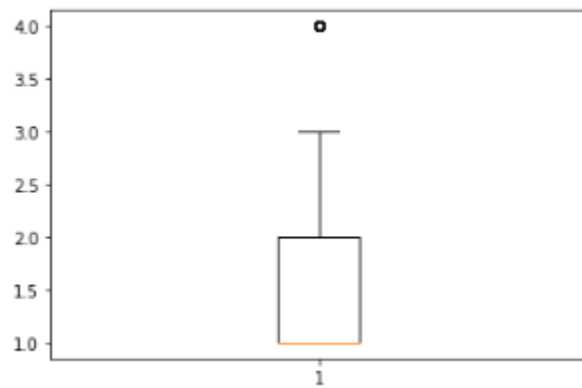
	Age	Balance	CreditScore	CustomerId	EstimatedSalary	Exited	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
...	...	...	...	...	...	...	
9995	False	False	False	False	False	False	
9996	False	False	False	False	False	False	
9997	False	False	False	False	False	False	
9998	False	False	False	False	False	False	
9999	False	False	False	False	False	False	

	Gender	Geography	HasCrCard	IsActiveMember	NumOfProducts	RowNumber	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
...	...	...	...	...	...	...	
9995	False	False	False	False	False	False	
9996	False	False	False	False	False	False	
9997	False	False	False	False	False	False	
9998	False	False	False	False	False	False	
9999	False	False	False	False	False	False	

	Surname	Tenure
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...	...	...
9995	False	False
9996	False	False
9997	False	False
9998	False	False
9999	False	False

```
plt.boxplot(df["NumOfProducts"])
```

```
plt.show()
```



```
np.where(df.Age>42,42, df.Age)
```

```
Out[16]: array([42, 41, 42, ..., 36, 42, 28], dtype=int64)
```

```
print(df['Age'].skew())
```

```
1.0113202630234552
```

```
print(df['Age'].quantile(0.25))
```

```
print(df['Age'].quantile(0.75))
```

```
df['Age'] = np.where(df['Age'] > 39, 41, df['Age'])
```

```
df.describe()
```

```
32.0
41.0
```

```
Out[22]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	1
mean	5000.50000	1.569094e+07	650.528800	35.788600	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.230881
std	2886.89568	7.193619e+04	96.653299	5.659409	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818
min	1.00000	1.559570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000
75%	7500.25000	1.575323e+07	718.000000	41.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500
max	10000.00000	1.581569e+07	850.000000	41.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000

## 7.categorical column

```
df["CustomerId"].value_counts()
```

```
Out[23]: 15812607    1
         15741078    1
         15635776    1
         15740223    1
         15738174    1
         ..
         15743714    1
         15639265    1
         15641312    1
         15684319    1
         15695872    1
         Name: CustomerId, Length: 10000, dtype: int64
```

```
df.dtypes
```

```
Out[27]: RowNumber      int64
         CustomerId     int64
         Surname        category
         CreditScore     int64
         Geography       object
         Gender          object
         Age             int64
         Tenure          int64
         Balance         float64
         NumOfProducts   int64
         HasCrCard       int64
         IsActiveMember  int64
         EstimatedSalary float64
         Exited          int64
         dtype: object
```

```
df["Age"].value_counts().sort_index()
```

```
Out[32]: 18    22
         19    27
         20    40
         21    53
         22    84
         23    99
         24   132
         25   154
         26   200
         27   209
         28   273
         29   348
         30   327
         31   404
         32   418
         33   442
         34   447
         35   474
         36   456
         37   478
         38   477
         39   423
         41  4013
         Name: Age, dtype: int64
```

```
df_categorical = df[categorical_columns]
```

```
df_categorical.head()
```

```
Out[35]:
```

	Geography	Gender
0	France	Female
1	Spain	Female
2	France	Female
3	France	Female
4	Spain	Female

```
pd.get_dummies(df, columns=["Age"]).head()
```

```
Out[36]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Tenure	Balance	NumOfProducts	HasCrCard	...	Age_31	Age_32	Age_33	Age_34
0	1	15634602	Hargrave	619	France	Female	2	0.00	1	1	...	0	0	0	0
1	2	15647311	Hill	608	Spain	Female	1	83807.86	1	0	...	0	0	0	0
2	3	15619304	Onio	502	France	Female	8	150660.80	3	1	...	0	0	0	0
3	4	15701354	Boni	699	France	Female	1	0.00	2	0	...	0	0	0	0
4	5	15737888	Mitchell	850	Spain	Female	2	125510.82	1	1	...	0	0	0	0

5 rows x 36 columns

## 8.split the data into dependent and independent variables

```
print(df.size)
```

```
140000
```

```
X = df.iloc[:, :-1].values
```

```
print(X)
```

```
[[1 15634602 'Hargrave' ... 1 1 101348.88]
 [2 15647311 'Hill' ... 0 1 112542.58]
 [3 15619304 'Onio' ... 1 0 113931.57]
 ...
 [9998 15584532 'Liu' ... 0 1 42085.58]
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]
 [10000 15628319 'Walker' ... 1 0 38190.78]]
```

```
Y = df.iloc[:, -1].values
```

```
print(Y)
```

```
[1 0 1 ... 1 1 0]
```

## 9.minmaxscaler

```
from sklearn.preprocessing import MinMaxScaler
```

```
df
```

```
scaler = MinMaxScaler()
```

```
print(scaler.fit(df))
```

```
MinMaxScaler(copy=True, feature_range=(0, 1))
```

## 10.train –split data

```
import pandas as pd
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split
```

```
df=pd.read_csv("C:\\Users\\PC\\Desktop\\Churn_Modelling.csv")
```

```
df.head()
```

```
Out[77]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15847311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159860.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

```
y= df.Tenure
```

```
y.head()
```

```
Out[78]: 0    2
         1    1
         2    8
         3    1
         4    2
         Name: Tenure, dtype: int64
```

```
x=df.drop('Tenure',axis=1)
```

```
x.head()
```

Out[80]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	83807.88	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	150060.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	0.00	2	0	0	93626.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	125510.82	1	1	1	79084.10	0

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

```
x_train.shape
```

Out[82]: (8000, 13)

```
y_train.shape
```

Out[83]: (8000,)

```
x_test.shape
```

Out[84]: (2000, 13)

```
y_test.shape
```

Out[85]: (2000,)