

### ASSESSMENT 3

|                     |              |
|---------------------|--------------|
| ASSESSMENT DATE     | 06-10-2022   |
| STUDENT NAME        | Saruthi.M    |
| STUDENT ROLL NUMBER | 713119205008 |
| MAXIMUM MARKS       | 2 Marks      |

```
import pandas as pd
import numpy as np
```

#### #1. Download the dataset

```
from google.colab import files
uploaded = files.upload()
```

#### #2. Load the dataset into the tool.

```
df=pd.read_csv('abalone - abalone.csv')
df.head()
```

|   | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings | Age  |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|------|
| 0 | M   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.150        | 15    | 16.5 |
| 1 | M   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.070        | 7     | 8.5  |
| 2 | F   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.210        | 9     | 10.5 |
| 3 | M   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.155        | 10    | 11.5 |
| 4 | I   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.055        | 7     | 8.5  |

```
0s
df['age'] = df['Rings']+1.5
df.drop('Rings', axis = 1, inplace = True)
df.age
```

```
[6] 0      16.5
    1       8.5
    2      10.5
    3      11.5
    4       8.5
    ...
   4172    12.5
   4173    11.5
   4174    10.5
   4175    11.5
   4176    13.5
    Name: age, Length: 4177, dtype: float64
```

```
df.shape
```

```
(4177, 10)
```

```
df.info()
```

```
[8] 0  Sex      4177 non-null  object
    1  Length  4177 non-null  float64
    2  Diameter 4177 non-null  float64
    3  Height  4177 non-null  float64
    4  Whole weight 4177 non-null float64
    5  Shucked weight 4177 non-null float64
    6  Viscera weight 4177 non-null float64
    7  Shell weight 4177 non-null float64
    8  Unnamed: 9  4177 non-null float64
    9  age      4177 non-null  float64
    dtypes: float64(9), object(1)
    memory usage: 326.5+ KB
```

```
# 3. Perform Below Visualizations.
```

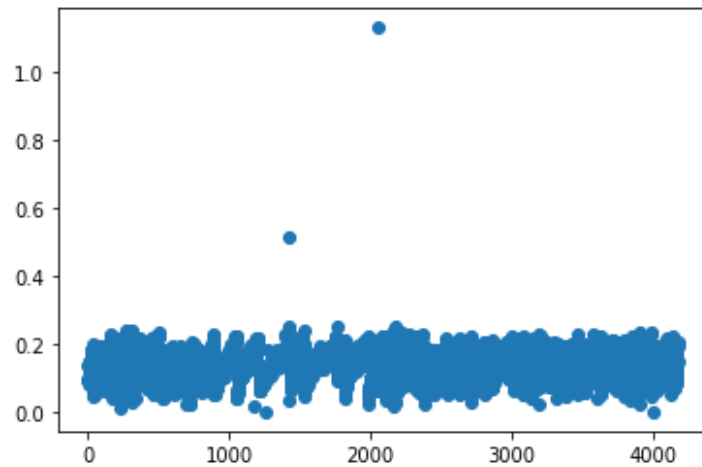
```
#univariate analysis
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```


```
plt.scatter(df.index,df['Height'])
```

```
plt.show()
```

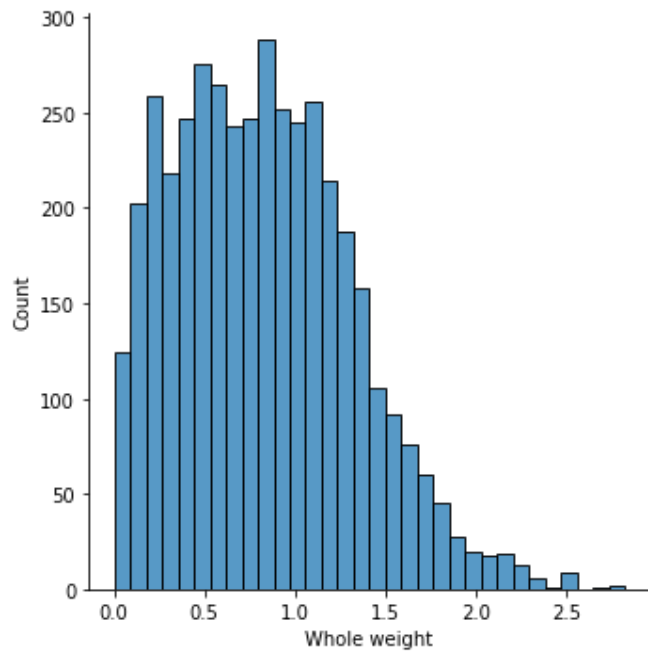


```
plt.hist(df['Diameter'])
```

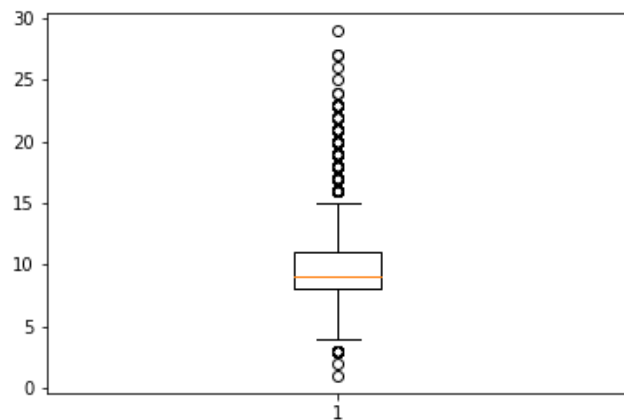
```
(array([ 13.,  66., 180., 344., 513., 812., 1017., 934., 275.,
        23.]),
 array([0.055 , 0.1145, 0.174 , 0.2335, 0.293 , 0.3525, 0.412 , 0.4715,
        0.531 , 0.5905, 0.65  ]),
 <a list of 10 Patch objects>)
```



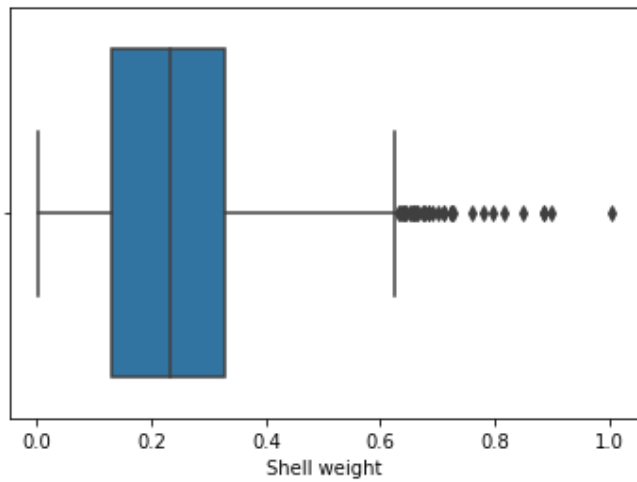
```
sns.displot(df['Whole weight'])
<seaborn.axisgrid.FacetGrid at 0x7f400117f050>
```



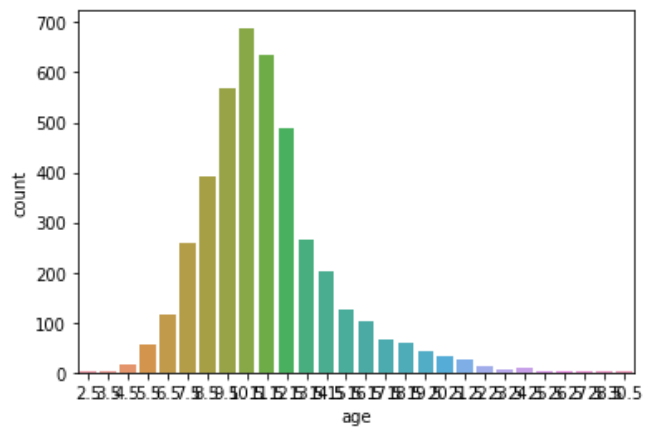
```
plt.boxplot(df['age'])
{'whiskers': [<matplotlib.lines.Line2D at 0x7f3ffe636e50>,
<matplotlib.lines.Line2D at 0x7f3ffe636b50>],
'caps': [<matplotlib.lines.Line2D at 0x7f3ffe6384d0>,
<matplotlib.lines.Line2D at 0x7f3ffe638b50>],
'boxes': [<matplotlib.lines.Line2D at 0x7f3ffe6368d0>],
'medians': [<matplotlib.lines.Line2D at 0x7f3ffe5dff90>],
'fliers': [<matplotlib.lines.Line2D at 0x7f3ffe5df790>],
'means': []}
```



```
sns.boxplot(df['Shell weight'])
```

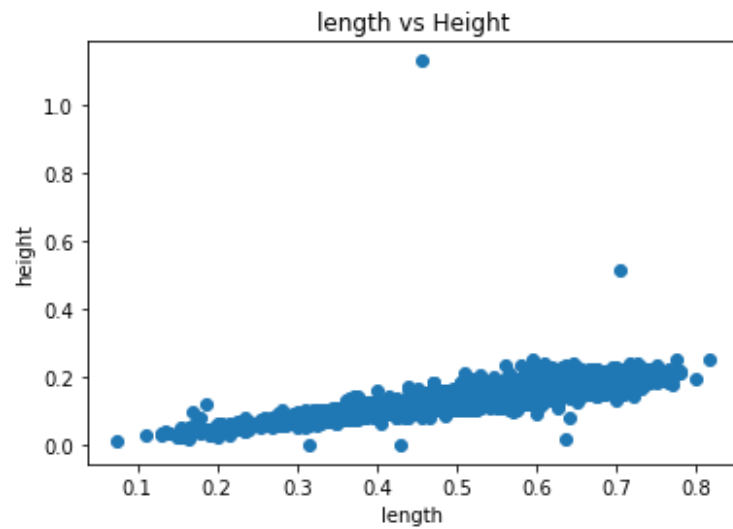


```
sns.countplot(df['age'])
```



### #Bivariate analysis

```
plt.scatter(df.Length, df.Height)
plt.title('length vs Height')
plt.xlabel('length')
plt.ylabel('height')
```



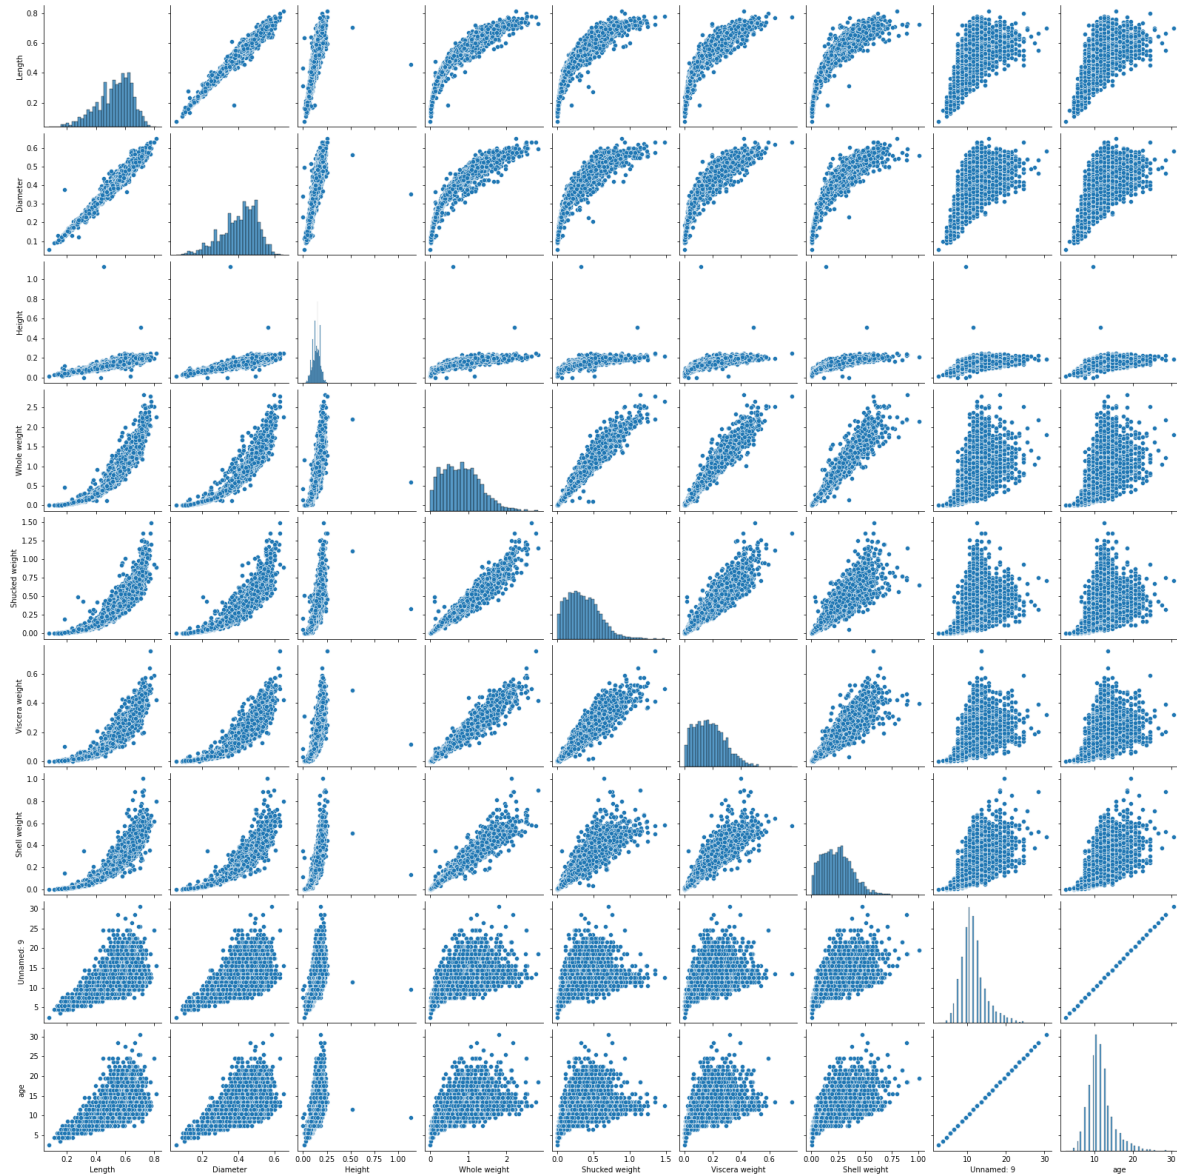
```
df.corr()
```

|                | Length   | Diameter | Height   | Whole weight | Shucked weight | Viscera weight | Shell weight | Unnamed: 9 | age      |
|----------------|----------|----------|----------|--------------|----------------|----------------|--------------|------------|----------|
| Length         | 1.000000 | 0.986812 | 0.827554 | 0.925261     | 0.897914       | 0.903018       | 0.897706     | 0.556720   | 0.556720 |
| Diameter       | 0.986812 | 1.000000 | 0.833684 | 0.925452     | 0.893162       | 0.899724       | 0.905330     | 0.574660   | 0.574660 |
| Height         | 0.827554 | 0.833684 | 1.000000 | 0.819221     | 0.774972       | 0.798319       | 0.817338     | 0.557467   | 0.557467 |
| Whole weight   | 0.925261 | 0.925452 | 0.819221 | 1.000000     | 0.969405       | 0.966375       | 0.955355     | 0.540390   | 0.540390 |
| Shucked weight | 0.897914 | 0.893162 | 0.774972 | 0.969405     | 1.000000       | 0.931961       | 0.882617     | 0.420884   | 0.420884 |
| Viscera weight | 0.903018 | 0.899724 | 0.798319 | 0.966375     | 0.931961       | 1.000000       | 0.907656     | 0.503819   | 0.503819 |
| Shell weight   | 0.897706 | 0.905330 | 0.817338 | 0.955355     | 0.882617       | 0.907656       | 1.000000     | 0.627574   | 0.627574 |

```
#multivariate analysis
```

```
sns.pairplot(df)
```

```
plt.show
```



# 4.Perform descriptive statistics on the dataset.

```
df.describe()
```

|       | Length      | Diameter    | Height      | Whole weight | Shucked weight | Viscera weight | Shell weight | Unnamed: 9  | age         |
|-------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|-------------|-------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000  | 4177.000000    | 4177.000000    | 4177.000000  | 4177.000000 | 4177.000000 |
| mean  | 0.523992    | 0.407881    | 0.139516    | 0.828742     | 0.359367       | 0.180594       | 0.238831     | 11.433684   | 11.433684   |
| std   | 0.120093    | 0.099240    | 0.041827    | 0.490389     | 0.221963       | 0.109614       | 0.139203     | 3.224169    | 3.224169    |
| min   | 0.075000    | 0.055000    | 0.000000    | 0.002000     | 0.001000       | 0.000500       | 0.001500     | 2.500000    | 2.500000    |
| 25%   | 0.450000    | 0.350000    | 0.115000    | 0.441500     | 0.186000       | 0.093500       | 0.130000     | 9.500000    | 9.500000    |
| 50%   | 0.545000    | 0.425000    | 0.140000    | 0.799500     | 0.336000       | 0.171000       | 0.234000     | 10.500000   | 10.500000   |

#5.Check for Missing values and deal with them.

```
df.info()
```

```

Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Sex                   4177 non-null   object
 1   Length                4177 non-null   float64
 2   Diameter              4177 non-null   float64
 3   Height               4177 non-null   float64
 4   Whole weight         4177 non-null   float64
 5   Shucked weight       4177 non-null   float64
 6   Viscera weight        4177 non-null   float64
 7   Shell weight         4177 non-null   float64
 8   Unnamed: 9           4177 non-null   float64
 9   age                  4177 non-null   float64
dtypes: float64(9), object(1)
memory usage: 326.5+ KB

```

```
df.isnull()
```

[illegible]

```
df.notnull()
```

[illegible]



```
df.fillna(0)
```

|     | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Unnamed: 9 | age  |
|-----|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------------|------|
| 0   | M   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.1500       | 16.5       | 16.5 |
| 1   | M   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.0700       | 8.5        | 8.5  |
| 2   | F   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.2100       | 10.5       | 10.5 |
| 3   | M   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.1550       | 11.5       | 11.5 |
| 4   | I   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.0550       | 8.5        | 8.5  |
| ... | ... | ...    | ...      | ...    | ...          | ...            | ...            | ...          | ...        | ...  |

```
df['Length'].fillna('No Length',inplace=True)
df.head()
```

|   | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Unnamed: 9 | age  |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|------------|------|
| 0 | M   | 0.455  | 0.365    | 0.095  | 0.5140       | 0.2245         | 0.1010         | 0.150        | 16.5       | 16.5 |
| 1 | M   | 0.350  | 0.265    | 0.090  | 0.2255       | 0.0995         | 0.0485         | 0.070        | 8.5        | 8.5  |
| 2 | F   | 0.530  | 0.420    | 0.135  | 0.6770       | 0.2565         | 0.1415         | 0.210        | 10.5       | 10.5 |
| 3 | M   | 0.440  | 0.365    | 0.125  | 0.5160       | 0.2155         | 0.1140         | 0.155        | 11.5       | 11.5 |
| 4 | I   | 0.330  | 0.255    | 0.080  | 0.2050       | 0.0895         | 0.0395         | 0.055        | 8.5        | 8.5  |

```
df.drop('Shucked weight',axis=1,inplace=True)
df.tail()
```

|      | Sex | Length | Diameter | Height | Whole weight | Viscera weight | Shell weight | Unnamed: 9 | age  |
|------|-----|--------|----------|--------|--------------|----------------|--------------|------------|------|
| 4172 | F   | 0.565  | 0.450    | 0.165  | 0.8870       | 0.2390         | 0.2490       | 12.5       | 12.5 |
| 4173 | M   | 0.590  | 0.440    | 0.135  | 0.9660       | 0.2145         | 0.2605       | 11.5       | 11.5 |
| 4174 | M   | 0.600  | 0.475    | 0.205  | 1.1760       | 0.2875         | 0.3080       | 10.5       | 10.5 |
| 4175 | F   | 0.625  | 0.485    | 0.150  | 1.0945       | 0.2610         | 0.2960       | 11.5       | 11.5 |
| 4176 | M   | 0.710  | 0.555    | 0.195  | 1.9485       | 0.3765         | 0.4950       | 13.5       | 13.5 |

```
print(df.isnull().sum())
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Viscera weight 0
Shell weight 0
Unnamed: 9   0
age          0
dtype: int64
```

**#6.Find the outliers and replace them outliers**

```
Q1=df.quantile(0.25)
```

```
Q3=df.quantile(0.75)
IQR=Q3-Q1
print(IQR)
```

```
Length      0.1650
Diameter     0.1300
Height      0.0500
Whole weight 0.7115
Viscera weight 0.1595
Shell weight 0.1990
Rings       3.0000
Age         3.0000
dtype: float64
```

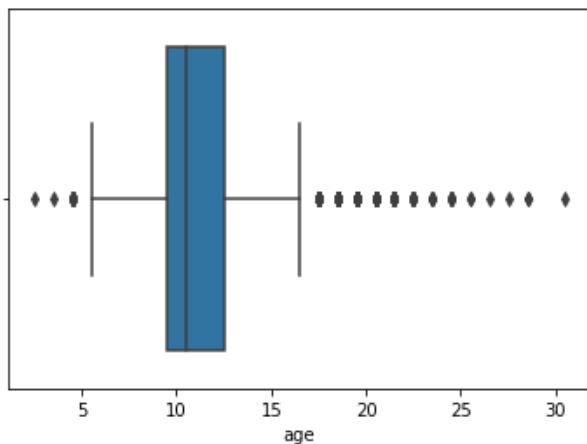
```
print(df<(Q1-1.5*IQR))
(df>(Q3+1.5*IQR))
```

```

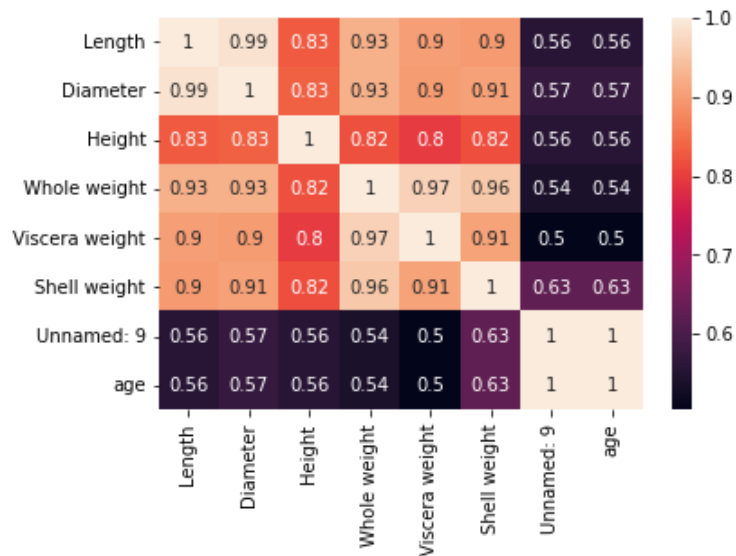
└─>   Age  Diameter  Height  Length  Rings  Sex  Shell weight  \
0     False      False   False   False  False  False      False
1     False      False   False   False  False  False      False
2     False      False   False   False  False  False      False
3     False      False   False   False  False  False      False
4     False      False   False   False  False  False      False
...     ...      ...     ...     ...     ...     ...      ...
4172  False      False   False   False  False  False      False
4173  False      False   False   False  False  False      False
4174  False      False   False   False  False  False      False
4175  False      False   False   False  False  False      False
4176  False      False   False   False  False  False      False

      Viscera weight  Whole weight
0                False          False
1                False          False
2                False          False
3                False          False
4                False          False
```

```
sns.boxplot(df.Rings)
```



```
sns.heatmap(df.corr(),annot=True)
```



```
np.where(df.age>7,7,df.age)
```

**output**

```
array([7, 7, 7, ..., 7, 7, 7])
```

```
print(df['Height'].quantile(0.25))
```

```
print(df['Height'].quantile(0.75))
```

```
df['Height']=np.where(df['Height']>0.090 ,0.125,df['Height'])
```

```
df.describe()
```

|       | Length      | Diameter    | Height      | Whole weight | Viscera weight | Shell weight | Unnamed: 9  | age         |
|-------|-------------|-------------|-------------|--------------|----------------|--------------|-------------|-------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000  | 4177.000000    | 4177.000000  | 4177.000000 | 4177.000000 |
| mean  | 0.523992    | 0.407881    | 0.118327    | 0.828742     | 0.180594       | 0.238831     | 11.433684   | 11.433684   |
| std   | 0.120093    | 0.099240    | 0.018405    | 0.490389     | 0.109614       | 0.139203     | 3.224169    | 3.224169    |
| min   | 0.075000    | 0.055000    | 0.000000    | 0.002000     | 0.000500       | 0.001500     | 2.500000    | 2.500000    |
| 25%   | 0.450000    | 0.350000    | 0.125000    | 0.441500     | 0.093500       | 0.130000     | 9.500000    | 9.500000    |
| 50%   | 0.545000    | 0.425000    | 0.125000    | 0.799500     | 0.171000       | 0.234000     | 10.500000   | 10.500000   |
| 75%   | 0.615000    | 0.480000    | 0.125000    | 1.153000     | 0.253000       | 0.329000     | 12.500000   | 12.500000   |

```
# 7.Check for Categorical columns and perform encoding.
```

```
df['Diameter'].value_counts()
```

---

```
0.450    139
0.475    120
0.400    111
0.500    110
0.470    100
...
0.610     1
0.650     1
0.620     1
0.095     1
0.615     1
Name: Diameter, Length: 111, dtype: int64
```

```
df.dtypes
```

```
Sex          object
Length       float64
Diameter     float64
Height       float64
Whole weight  float64
Viscera weight float64
Shell weight  float64
Unnamed: 9    float64
age          float64
dtype: object
```

```
df['Whole weight'].value_counts().sort_index()
```

---

```
0.0020    1
0.0080    1
0.0105    1
0.0130    1
0.0140    1
..
2.5500    1
2.5550    1
2.6570    1
2.7795    1
2.8255    1
Name: Whole weight, Length: 2429, dtype: int64
```

```
pd.get_dummies(df,columns=['Whole weight']).tail()
```

|      | Sex | Length | Diameter | Height | Viscera weight | Shell weight | Unnamed: 9 | age  | Whole weight_0.002 | Whole weight_0.008 | ... | Whole weight_2.505 | Whole weight_2.506 |
|------|-----|--------|----------|--------|----------------|--------------|------------|------|--------------------|--------------------|-----|--------------------|--------------------|
| 4172 | F   | 0.565  | 0.450    | 0.125  | 0.2390         | 0.2490       | 12.5       | 12.5 | 0                  | 0                  | ... | 0                  | 0                  |
| 4173 | M   | 0.590  | 0.440    | 0.125  | 0.2145         | 0.2605       | 11.5       | 11.5 | 0                  | 0                  | ... | 0                  | 0                  |
| 4174 | M   | 0.600  | 0.475    | 0.125  | 0.2875         | 0.3080       | 10.5       | 10.5 | 0                  | 0                  | ... | 0                  | 0                  |
| 4175 | F   | 0.625  | 0.485    | 0.125  | 0.2610         | 0.2960       | 11.5       | 11.5 | 0                  | 0                  | ... | 0                  | 0                  |
| 4176 | M   | 0.710  | 0.555    | 0.125  | 0.3765         | 0.4950       | 13.5       | 13.5 | 0                  | 0                  | ... | 0                  | 0                  |

5 rows × 2437 columns

```
from sklearn.preprocessing import OneHotEncoder
```

```
one_encde= OneHotEncoder(sparse=False)
```

```
encoded_arr=one_encde.fit_transform(df[['Length','Diameter','Height','Viscera weight']])
```

```
encoded_arr
```

```
array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]])
```

#8. Split the data into dependent and independent variables.

```
x=df.iloc[:,1:3]
```

```
x
```

↗

|      | Length | Diameter |
|------|--------|----------|
| 0    | 0.455  | 0.365    |
| 1    | 0.350  | 0.265    |
| 2    | 0.530  | 0.420    |
| 3    | 0.440  | 0.365    |
| 4    | 0.330  | 0.255    |
| ...  | ...    | ...      |
| 4172 | 0.565  | 0.450    |
| 4173 | 0.590  | 0.440    |
| 4174 | 0.600  | 0.475    |

```
y=df.iloc[:,1:4]
```

```
y
```

|      | Length | Diameter | Height |
|------|--------|----------|--------|
| 0    | 0.455  | 0.365    | 0.095  |
| 1    | 0.350  | 0.265    | 0.090  |
| 2    | 0.530  | 0.420    | 0.135  |
| 3    | 0.440  | 0.365    | 0.125  |
| 4    | 0.330  | 0.255    | 0.080  |
| ...  | ...    | ...      | ...    |
| 4172 | 0.565  | 0.450    | 0.165  |
| 4173 | 0.590  | 0.440    | 0.135  |

### #9. Scale the independent variables.

```
from sklearn.preprocessing import MinMaxScaler
model=MinMaxScaler()
scaled_x=pd.DataFrame(model.fit_transform(x),columns=x.columns)
scaled_x.head()
```

|   | Length   | Diameter |
|---|----------|----------|
| 0 | 0.513514 | 0.521008 |
| 1 | 0.371622 | 0.352941 |
| 2 | 0.614865 | 0.613445 |
| 3 | 0.493243 | 0.521008 |
| 4 | 0.344595 | 0.336134 |

### #10. Split the data into training and testing

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
x_train.shape
```

```
(3341, 2)
```

```
x_test.shape
```

```
(836, 2)
```

```
y_train.shape
```

```
(3341, 3)
```

```
y_test.shape
```

```
(836, 3)
```

### #11. Build the Model

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(x,y)
```

**output**

```
LinearRegression()
```

### #12. Train the Model

```
model.fit(x_train, y_train)
```

```
LinearRegression()
```

**output**

```
LinearRegression()
```

### #13. Test the Model

```
pred1=model.predict(x_train)
```

```
pred1
```

```
array([[0.18      , 0.135      , 0.08321098],
       [0.215      , 0.15       , 0.08626508],
       [0.66       , 0.53       , 0.13284892],
       ...,
       [0.595      , 0.45       , 0.12504125],
       [0.625      , 0.49       , 0.12877082],
       [0.41       , 0.325      , 0.10702569]])
```

```
predictions=model.predict(x_test)
```

```
predictions
```

```
pred=model.predict(x_test)
```

```
pred
```

```
array([[0.18      , 0.135      , 0.08321098],
       [0.215      , 0.15      , 0.08626508],
       [0.66      , 0.53      , 0.13284892],
       ...,
       [0.595      , 0.45      , 0.12504125],
       [0.625      , 0.49      , 0.12877082],
       [0.41      , 0.325      , 0.10702569]])
```

```
pred=model.predict(x_test)
```

```
pred
```

```
y_pred=(x_test)
```

```
y_pred
```

|      | Length | Diameter |
|------|--------|----------|
| 668  | 0.550  | 0.425    |
| 1580 | 0.500  | 0.400    |
| 3784 | 0.620  | 0.480    |
| 463  | 0.220  | 0.165    |
| 2615 | 0.645  | 0.500    |

```
y_test
```

|      | Length | Diameter | Height |
|------|--------|----------|--------|
| 668  | 0.550  | 0.425    | 0.125  |
| 1580 | 0.500  | 0.400    | 0.125  |
| 3784 | 0.620  | 0.480    | 0.125  |
| 463  | 0.220  | 0.165    | 0.055  |
| 2615 | 0.645  | 0.500    | 0.125  |
| ...  | ...    | ...      | ...    |
| 575  | 0.610  | 0.475    | 0.125  |



```
length=pd.DataFrame({'Actual_y_value':[pred1],'predicted_y_value':[pred]})
length
```

|   | Actual_y_value                                      | predicted_y_value                                   |
|---|---|---|
| 0 | [[0.180000000000000055, 0.135000000000000023, 0.... | [[0.180000000000000055, 0.135000000000000023, 0.... |

```
from sklearn import metrics
#Mean Absolute Error (MAE)
metrics.mean_absolute_error(y_test, predictions)
0.0038991947837602914
#Mean Squared Error (MSE)
metrics.mean_squared_error(y_test, predictions)
7.655875085238909e-05
#Root Mean Squared Error (RMSE)
np.sqrt(metrics.mean_squared_error(y_test, predictions))
0.008749785760370884
#14. Measure the performance using Metrics.
from sklearn.metrics import accuracy_score
accuracy_score=(y_test,y_pred)
accuracy_score
```

| (    | Length | Diameter | Height |
|------|--------|----------|--------|
| 668  | 0.550  | 0.425    | 0.125  |
| 1580 | 0.500  | 0.400    | 0.125  |
| 3784 | 0.620  | 0.480    | 0.125  |
| 463  | 0.220  | 0.165    | 0.055  |
| 2615 | 0.645  | 0.500    | 0.125  |
| ...  | ...    | ...      | ...    |
| 575  | 0.610  | 0.475    | 0.125  |
| 3231 | 0.410  | 0.325    | 0.125  |
| 1084 | 0.445  | 0.345    | 0.125  |
| 290  | 0.540  | 0.435    | 0.125  |
| 2713 | 0.250  | 0.175    | 0.060  |

```
from sklearn.metrics import classification_report
classification_report(y_test,y_pred)
classification_report
```

```
<function sklearn.metrics._classification.classification_report(y_true,
y_pred, *, labels=None, target_names=None, sample_weight=None, digits=2,
output_dict=False, zero_division='warn')>
```

